

Question2

May 7, 2025

0.0.1 Name:Ying Zhou

0.0.2 Question 2A

I extracted solar_data.tar.gz, which contains 10 .txt files from 2014 to 2023. Each file starts with metadata, followed by daily solar data in table format. I checked all 10 files and confirmed they have the same structure: date, radio flux, sunspot data, solar field, X-ray flares (C/M/X), and optical flares (S, 1, 2, 3). Fields are whitespace-separated and consistent across files.

```
In [23]: import os
import tarfile

!cp /public/bmort/python/solar_data.tar.gz .
with tarfile.open("solar_data.tar.gz", "r:gz") as tar:
    tar.extractall("solar_data")

data_path = "solar_data/solar_data"
txt_files = sorted([f for f in os.listdir(data_path) if f.endswith('.txt')])

for fname in txt_files:
    print(f"\n Examining file: {fname}")
    file_path = os.path.join(data_path, fname)
    with open(file_path, 'r') as f:
        line_count = 0
        for line in f:
            if not line.startswith('#') and not line.startswith(':') and line.strip():
                print(line.strip())
                line_count += 1
            if line_count >= 4: # Only show 4 data lines per file
                break
```

```
Examining file: 2014_DSD.txt
2014 01 01 160 106 610 1 -999 C1.1 3 1 0 7 0 1 0
2014 01 02 161 133 1410 1 -999 C1.5 6 2 0 5 1 0 0
2014 01 03 182 162 1570 0 -999 C1.2 9 2 0 14 0 0 0
2014 01 04 215 178 1710 1 -999 C1.7 4 3 0 10 0 1 0
```

```
Examining file: 2015_DSD.txt
```

2015	01	01	138	101	870	1	-999	B4.9	3	0	0	9	0	0	0
2015	01	02	146	113	1250	1	-999	B4.9	3	0	0	9	0	0	0
2015	01	03	149	122	1300	0	-999	B5.7	10	1	0	13	1	0	0
2015	01	04	150	124	1220	0	-999	B6.7	5	1	0	13	0	1	0

Examining file: 2016_DSD.txt

2016	01	01	98	40	140	2	-999	B3.4	3	1	0	2	0	0	0
2016	01	02	100	52	140	1	-999	B5.8	0	0	0	1	0	0	0
2016	01	03	102	50	170	0	-999	B3.1	0	0	0	1	0	0	0
2016	01	04	95	60	160	0	-999	B2.2	0	0	0	0	0	0	0

Examining file: 2017_DSD.txt

2017	01	01	73	0	0	1	-999	A5.7	0	0	0	0	0	0	0
2017	01	02	73	0	0	0	-999	A5.7	0	0	0	0	0	0	0
2017	01	03	73	11	0	1	-999	A5.5	0	0	0	0	0	0	0
2017	01	04	72	0	0	0	-999	A5.4	0	0	0	0	0	0	0

Examining file: 2018_DSD.txt

2018	01	01	69	0	0	0	-999	A3.5	0	0	0	0	0	0	0
2018	01	02	70	0	0	0	-999	A3.2	0	0	0	0	0	0	0
2018	01	03	71	0	0	0	-999	A3.1	0	0	0	0	0	0	0
2018	01	04	70	13	20	1	-999	A3.1	0	0	0	0	0	0	0

Examining file: 2019_DSD.txt

2019	01	01	72	13	10	1	-999	A1.0	0	0	0	0	0	0	0
2019	01	02	75	16	30	0	-999	A1.9	0	0	0	1	0	0	0
2019	01	03	73	16	90	0	-999	A2.3	0	0	0	0	0	0	0
2019	01	04	72	13	50	0	-999	A1.5	0	0	0	0	0	0	0

Examining file: 2020_DSD.txt

2020	01	01	72	0	0	0	-999	A8.0	0	0	0	0	0	0	0
2020	01	02	72	13	20	1	-999	A8.3	0	0	0	0	0	0	0
2020	01	03	71	13	20	0	-999	A8.6	0	0	0	0	0	0	0
2020	01	04	72	11	10	0	-999	A8.5	0	0	0	0	0	0	0

Examining file: 2021_DSD.txt

2021	01	01	80	23	190	0	-999	*	0	0	0	0	0	0	0
2021	01	02	82	22	120	0	-999	*	0	0	0	0	0	0	0
2021	01	03	80	0	0	0	-999	*	0	0	0	0	0	0	0
2021	01	04	78	0	0	0	-999	*	0	0	0	0	0	0	0

Examining file: 2022_DSD.txt

2022	01	01	94	52	440	1	-999	*	1	1	0	1	0	0	0
2022	01	02	89	25	340	0	-999	*	1	0	0	0	0	0	0
2022	01	03	84	12	140	0	-999	*	0	0	0	0	0	0	0
2022	01	04	86	12	30	1	-999	*	0	0	0	0	0	0	0

Examining file: 2023_DSD.txt

2023	01	01	153	94	1220	1	-999	*	9	0	0	12	0	0	0
2023	01	02	146	94	1100	0	-999	*	9	0	0	7	1	0	0
2023	01	03	149	89	930	0	-999	*	5	0	0	1	0	0	0
2023	01	04	151	86	550	0	-999	*	3	0	0	4	0	1	0

0.0.3 Question 2B

```
In [24]: import pandas as pd
import glob

data_files = sorted(glob.glob('./solar_data/solar_data/*.txt'))

col_names = [
    'Date', 'RadioFlux', 'SunspotNumber', 'SunspotArea', 'NewRegions',
    'SolarMeanField', 'XrayBkgd', 'Xray_C', 'Xray_M', 'Xray_X',
    'Optical_S', 'Optical_1', 'Optical_2', 'Optical_3'
]

df_list = []

for f in data_files:
    try:
        df = pd.read_csv(
            f,
            delim_whitespace=True,
            skiprows=13,
            comment='#',
            header=None
        )

        year = df[0].astype(str)
        month = df[1].astype(str).str.zfill(2)
        day = df[2].astype(str).str.zfill(2)
        date_str = year + '-' + month + '-' + day

        df['Date'] = date_str

        df.columns = ['Year', 'Month', 'Day', 'RadioFlux', 'SunspotNumber', 'SunspotArea',
                      'SolarMeanField', 'XrayBkgd', 'Xray_C', 'Xray_M', 'Xray_X',
                      'Optical_S', 'Optical_1', 'Optical_2', 'Optical_3', 'Date']

        df = df.drop(columns=['Year', 'Month', 'Day'])

        df_list.append(df)

    except Exception as e:
        print(f"Failed to load {f}: {e}")
```

```
df_all = pd.concat(df_list, ignore_index=True)
df_all.head()
```

```
Out[24]:
```

	RadioFlux	SunspotNumber	SunspotArea	NewRegions	SolarMeanField	XrayBkgd	\
0	160	106	610	1	-999	C1.1	
1	161	133	1410	1	-999	C1.5	
2	182	162	1570	0	-999	C1.2	
3	215	178	1710	1	-999	C1.7	
4	218	225	1790	0	-999	C1.8	

	Xray_C	Xray_M	Xray_X	Optical_S	Optical_1	Optical_2	Optical_3	\
0	3	1	0	7	0	1	0	
1	6	2	0	5	1	0	0	
2	9	2	0	14	0	0	0	
3	4	3	0	10	0	1	0	
4	3	0	0	14	0	0	0	

	Date
0	2014-01-01
1	2014-01-02
2	2014-01-03
3	2014-01-04
4	2014-01-05

```
In [25]: import numpy as np
```

```
numeric_cols = df_all.columns.drop(['Date', 'XrayBkgd']).tolist()
```

```
df_all[numeric_cols] = df_all[numeric_cols].astype(float)
```

```
df_all[numeric_cols] = df_all[numeric_cols].replace(-999, np.nan)
```

```
summary = df_all[numeric_cols].describe()
display(summary)
```

	RadioFlux	SunspotNumber	SunspotArea	NewRegions	SolarMeanField	\
count	3650.000000	3650.000000	3650.000000	3650.000000	0.0	
mean	100.773699	49.189863	344.112877	0.433973	NaN	
std	35.477323	51.333475	449.820622	0.735244	NaN	
min	-1.000000	0.000000	0.000000	0.000000	NaN	
25%	71.000000	0.000000	0.000000	0.000000	NaN	
50%	85.000000	31.000000	160.000000	0.000000	NaN	
75%	126.000000	81.000000	540.000000	1.000000	NaN	
max	343.000000	296.000000	3120.000000	6.000000	NaN	

	Xray_C	Xray_M	Xray_X	Optical_S	Optical_1	\
count	3650.000000	3650.000000	3650.000000	3650.000000	3650.000000	

mean	2.498904	0.261096	0.011781	3.291507	0.210959
std	3.902361	0.854465	0.112878	5.588125	0.629407
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	4.000000	0.000000	0.000000	4.000000	0.000000
max	30.000000	11.000000	2.000000	55.000000	7.000000

	Optical_2	Optical_3
count	3650.000000	3650.000000
mean	0.034247	0.003836
std	0.194976	0.061822
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	0.000000	0.000000
max	2.000000	1.000000

0.0.4 Question 2B Summary

Loaded all yearly solar data into a single Pandas DataFrame.

Data cleaning actions: - Combined "Year, Month, Day" into a single Date column. - Replaced -999 with NaN to handle missing values. - Converted numeric columns to "float" type, excluding non-numeric ones like "XrayBkgd".

Summary statistics: - The numerical columns have different magnitudes and ranges. - For example, "SunspotArea" ranges up to 3120, while "Xray_M", "Optical_2", etc., have mostly zeros and small values. - This shows the variables are not on the same scale. - Some columns show potential outliers, especially those with large max values far above the 75th percentile.

0.0.5 Question2C

```
In [26]: max_flux = df_all['RadioFlux'].max()
max_days = df_all[df_all['RadioFlux'] == max_flux][['Date', 'RadioFlux']]

min_flux = df_all['RadioFlux'].min()
min_days = df_all[df_all['RadioFlux'] == min_flux][['Date', 'RadioFlux']]

display(max_days)
display(min_days)
```

	Date	RadioFlux
3333	2023-02-17	343.0

	Date	RadioFlux
3271	2022-12-16	-1.0

2023-02-17 had the highest value for the radio flux 2022-12-16 had the lowest value for the radio flux.

0.0.6 Question2D

```
In [27]: df_all['TotalFlares'] = df_all[['Optical_S', 'Optical_1', 'Optical_2', 'Optical_3']].sum(axis=1)

max_flare_value = df_all['TotalFlares'].max()
max_flare_days = df_all[df_all['TotalFlares'] == max_flare_value][['Date', 'TotalFlares']]

no_flare_days_count = (df_all['TotalFlares'] == 0).sum()

print("Day(s) with highest solar flares:")
display(max_flare_days)

print(f" Number of days with no solar flares: {no_flare_days_count}")
```

Day(s) with highest solar flares:

	Date	TotalFlares
3327	2023-02-11	60.0

Number of days with no solar flares: 1823

Day(s) with the highest number of solar flares: 2023-02-11 had the highest number of solar flares, with a total of 60 flares observed. Number of days with no solar flares: There were 1,823 days with zero solar flares recorded.

0.0.7 Question 2E

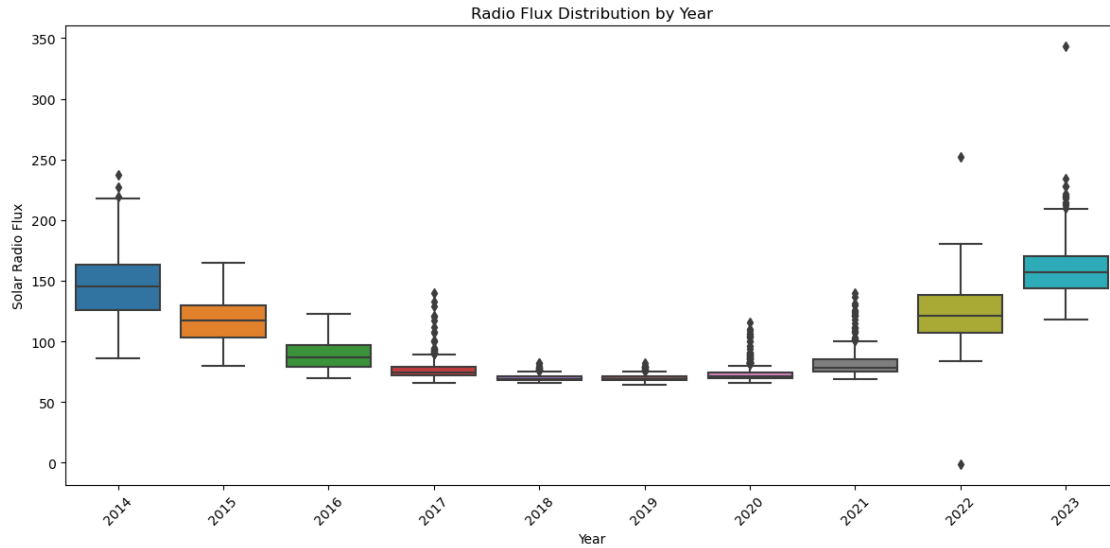
```
In [29]: import seaborn as sns
import matplotlib.pyplot as plt

df_all['Date'] = pd.to_datetime(df_all['Date'])

df_all['Year'] = df_all['Date'].dt.year

plt.figure(figsize=(12, 6))
sns.boxplot(x='Year', y='RadioFlux', data=df_all)

plt.title('Radio Flux Distribution by Year')
plt.xlabel('Year')
plt.ylabel('Solar Radio Flux')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



0.0.8 Question2F

```
In [30]: from sklearn.model_selection import train_test_split

train_df, test_df = train_test_split(df_all, test_size=0.2, random_state=42)

print(f"Training set size: {len(train_df)}")
print(f"Testing set size: {len(test_df)}")
```

Training set size: 2920

Testing set size: 730

0.0.9 Question2G and H

```
In [36]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

features = ['SunspotNumber', 'Xray_C', 'Xray_M', 'Xray_X',
            'Optical_S', 'Optical_1', 'Optical_2', 'Optical_3']
target = 'RadioFlux'

train_clean = train_df.dropna(subset=features + [target])
test_clean = test_df.dropna(subset=features + [target])

X_train = train_clean[features]
y_train = train_clean[target]
X_test = test_clean[features]
y_test = test_clean[target]
```

```

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", round(mse, 2))
print("R-squared:", round(r2, 4))

```

Mean Squared Error: 173.82
R-squared: 0.8644

The model predicts the solar radio flux quite well, with an R^2 score of approximately 0.86, indicating that about 86% of the variability in solar radio flux can be explained by sunspot number and solar flare activity. The accuracy is good because solar activity indicators like sunspots and flares have a strong correlation with radio flux. However, the prediction is not perfect due to possible noise in the data or other influencing factors not included in the model.

0.0.10 Question2I

```

In [37]: import pandas as pd

input_data = pd.DataFrame([
    'SunspotNumber': 96,
    'Xray_C': 1,
    'Xray_M': 0,
    'Xray_X': 0,
    'Optical_S': 0,
    'Optical_1': 0,
    'Optical_2': 0,
    'Optical_3': 0
])

predicted_flux = model.predict(input_data)[0]
print(f"Predicted Solar Radio Flux: {predicted_flux:.2f}")

```

Predicted Solar Radio Flux: 122.72

The predicted solar radio flux for a day with 96 sunspots and a single C-class X-ray flare is approximately 122.72.

We are reasonably confident in this prediction because the model achieves a high R^2 score of approximately 0.86. However, the prediction is not perfect due to noise in the data and other influencing factors not captured in the model.